

# Distance Vector (DV) Algorithm

- **Distributed:** each node receives some information from one or more of its *directly attached* neighbors, performs a calculation, and then distributes the results of its calculation back to its neighbors
- **Iterative:** this process continues on until no more information is exchanged between neighbors.
- **Asynchronous:** it does not require all of the nodes to operate in lockstep with each other.

.

$d_x(y)$  Cost of the least-cost path from node  $x$  to node  $y$

## **Bellman-Ford Equation**

$$d_x(y) = \min_v \{ c(x, v) + d_v(y) \}$$

Min over all neighbors

Each node  $x$  begins with:

$D_x(y)$ : Estimate of the best cost from itself to any node  $y$

## Distance Vector from $x$

$$\mathbf{D}_x = [D_x(y): y \text{ in } N]$$

Information at each node:


- For each neighbor  $v$ , the cost  $c(x,v)$  from  $x$  to directly attached neighbor,  $v$
- Node  $x$ 's distance vector, that is,  $\mathbf{D}_x = [D_x(y): y \text{ in } N]$ , containing  $x$ 's estimate of its cost to all destinations,  $y$ , in  $N$
- The distance vectors of each of its neighbors, that is,  $\mathbf{D}_v = [D_v(y): y \text{ in } N]$  for each neighbor  $v$  of  $x$

If node  $x$  receives a new DV from a neighbor (only local exchanges), then update:

$$D_x(y) = \min_v \{c(x,v) + D_v(y)\} \quad \text{for each node } y \text{ in } N$$

If DV changes, then  $x$  sends update to neighbors.

```
1  Initialization:
2      for all destinations  $y$  in  $N$ :
3           $D_x(y) = c(x,y)$  /* if  $y$  is not a neighbor then  $c(x,y) = \infty$  */
4      for each neighbor  $w$ 
5           $D_w(y) = ?$  for all destinations  $y$  in  $N$ 
6      for each neighbor  $w$ 
7          send distance vector  $\mathbf{D}_x = [D_x(y): y \text{ in } N]$  to  $w$ 
8
9  loop
10     wait (until I see a link cost change to some neighbor  $w$  or
11           until I receive a distance vector from some neighbor  $w$ )
12
13     for each  $y$  in  $N$ :
14          $D_x(y) = \min_v \{c(x,v) + D_v(y)\}$ 
15
16     if  $D_x(y)$  changed for any destination  $y$ 
17         send distance vector  $\mathbf{D}_x = [D_x(y): y \text{ in } N]$  to all neighbors
18
19 forever
```



Node x table

		cost to		
		x	y	z
from	x	0	2	7
	y	$\infty$	$\infty$	$\infty$
	z	$\infty$	$\infty$	$\infty$

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	7	1	0

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	3	1	0

Node y table

		cost to		
		x	y	z
from	x	$\infty$	$\infty$	$\infty$
	y	2	0	1
	z	$\infty$	$\infty$	$\infty$

		cost to		
		x	y	z
from	x	0	2	7
	y	2	0	1
	z	7	1	0

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	3	1	0

Node z table


		cost to		
		x	y	z
from	x	$\infty$	$\infty$	$\infty$
	y	$\infty$	$\infty$	$\infty$
	z	7	1	0

		cost to		
		x	y	z
from	x	0	2	7
	y	2	0	1
	z	3	1	0


		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	3	1	0

.....> Time

# Issues



a.




b.


## a) DV is updated efficiently

- At time  $t_0$ , y detects the link-cost change (the cost has changed from 4 to 1), updates its distance vector, and informs its neighbors of this change since its distance vector has changed.
- At time  $t_1$ , z receives the update from y and updates its table. It computes a new least cost to x (it has decreased from a cost of 5 to a cost of 2) and sends its new distance vector to its neighbors.
- At time  $t_2$ , y receives z's update and updates its distance table. y's least costs do not change and hence y does not send any message to z. The algorithm comes to a quiescent state.

# Issues



a.



b.

## b) Slow Convergence (44 iterations)

1. Before the link cost changes,  $D_y(x) = 4$ ,  $D_y(z) = 1$ ,  $D_z(y) = 1$ , and  $D_z(x) = 5$ . At time  $t_0$ ,  $y$  detects the link-cost change (the cost has changed from 4 to 60).  $y$  computes its new minimum-cost path to  $x$  to have a cost of

$$D_y(x) = \min\{c(y,x) + D_x(x), c(y,z) + D_z(x)\} = \min\{60 + 0, 1 + 5\} = 6$$

Of course, with our global view of the network, we can see that this new cost via  $z$  is *wrong*. But the only information node  $y$  has is that its direct cost to  $x$  is 60 and that  $z$  has last told  $y$  that  $z$  could get to  $x$  with a cost of 5. So in order to get to  $x$ ,  $y$  would now route through  $z$ , fully expecting that  $z$  will be able to get to  $x$  with a cost of 5. As of  $t_1$  we have a **routing loop**—in order to get to  $x$ ,  $y$  routes through  $z$ , and  $z$  routes through  $y$ . A routing loop is like a black hole—a packet destined for  $x$  arriving at  $y$  or  $z$  as of  $t_1$  will bounce back and forth between these two nodes forever (or until the forwarding tables are changed).

2. Since node  $y$  has computed a new minimum cost to  $x$ , it informs  $z$  of its new distance vector at time  $t_1$ .
3. Sometime after  $t_1$ ,  $z$  receives  $y$ 's new distance vector, which indicates that  $y$ 's minimum cost to  $x$  is 6.  $z$  knows it can get to  $y$  with a cost of 1 and hence computes a new least cost to  $x$  of  $D_z(x) = \min\{50 + 0, 1 + 6\} = 7$ . Since  $z$ 's least cost to  $x$  has increased, it then informs  $y$  of its new distance vector at  $t_2$ .
4. In a similar manner, after receiving  $z$ 's new distance vector,  $y$  determines  $D_y(x) = 8$  and sends  $z$  its distance vector.  $z$  then determines  $D_z(x) = 9$  and sends  $y$  its distance vector, and so on.



## **Solution: Poisoned Reverse**

If z routes through y to get to x, then z will advertise to y that its distance to x is infinity