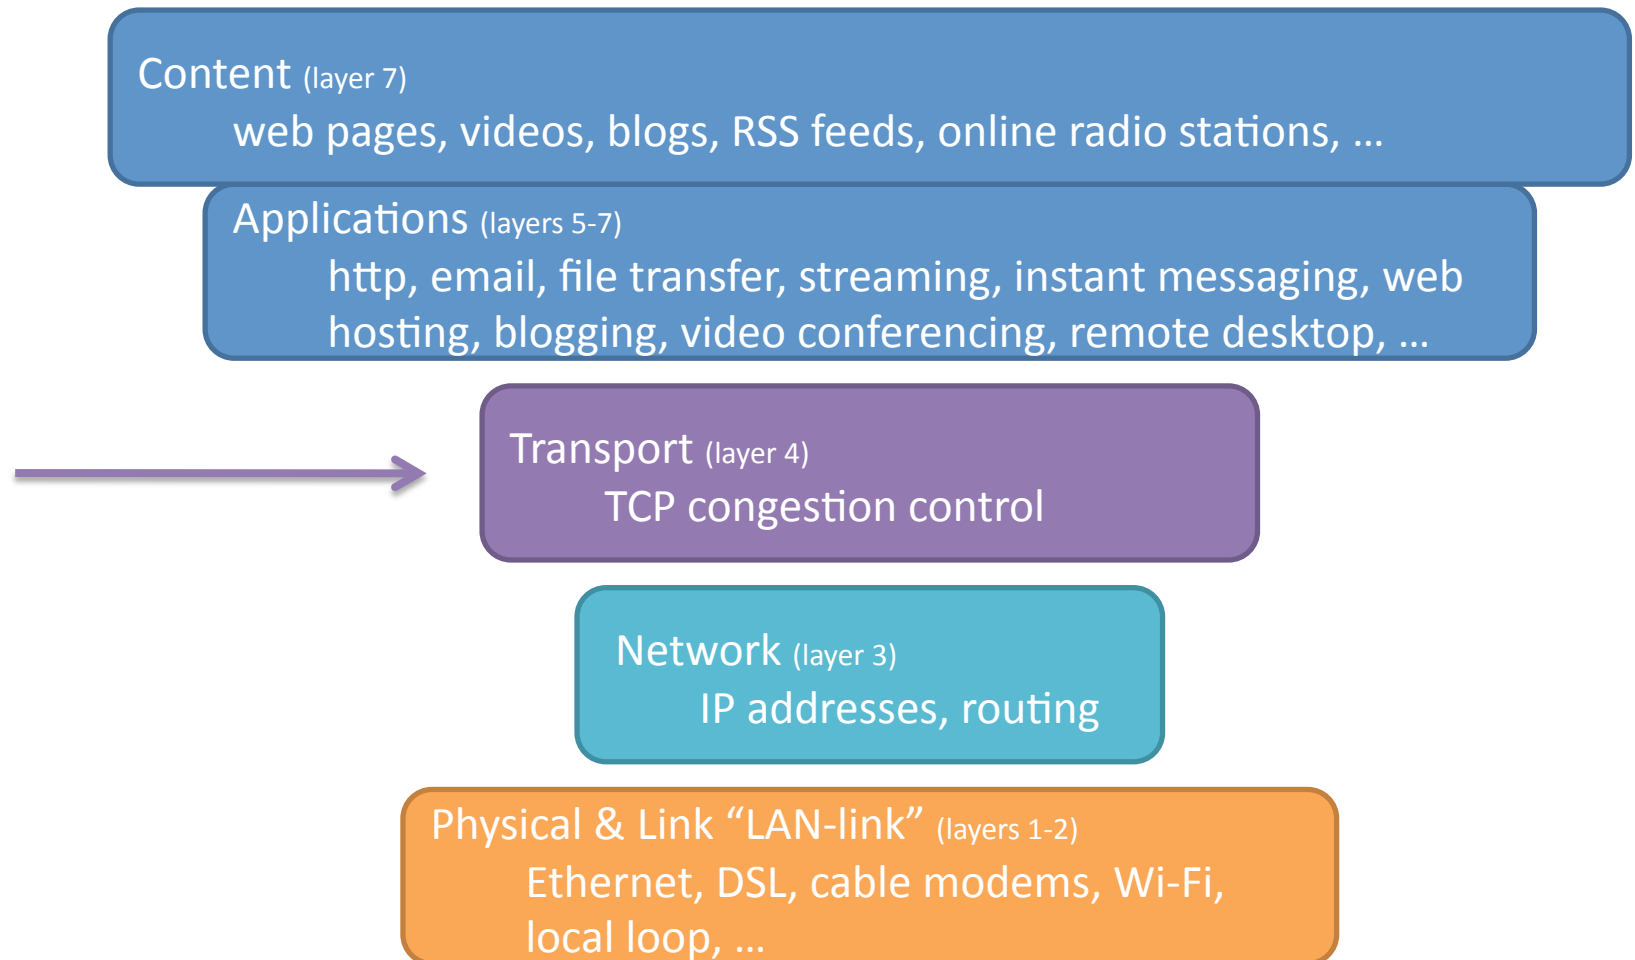


CS 232

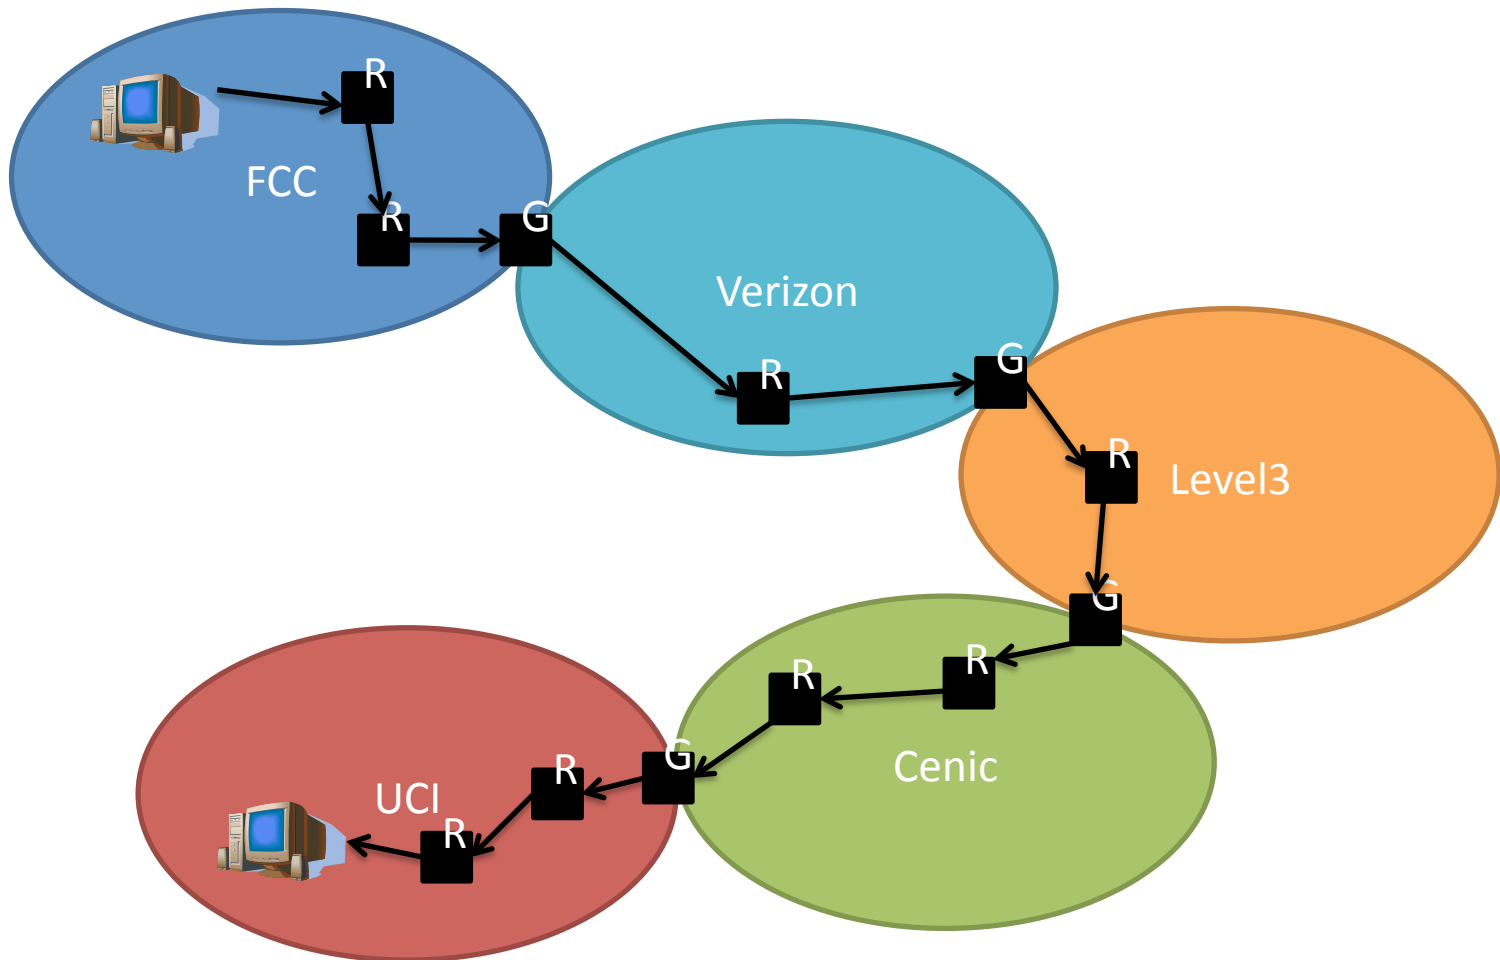
Transmission Control Protocol (TCP)

Some figures on these slides are reproduced from textbooks, and are provided under Fair Use solely to those enrolled in this course. The remainder of these slides are copyright Scott Jordan. Unauthorized reproduction or distribution (including posting on a website) of any portion of these slides is a violation of the UCI Student Code of Conduct and may constitute copyright infringement.

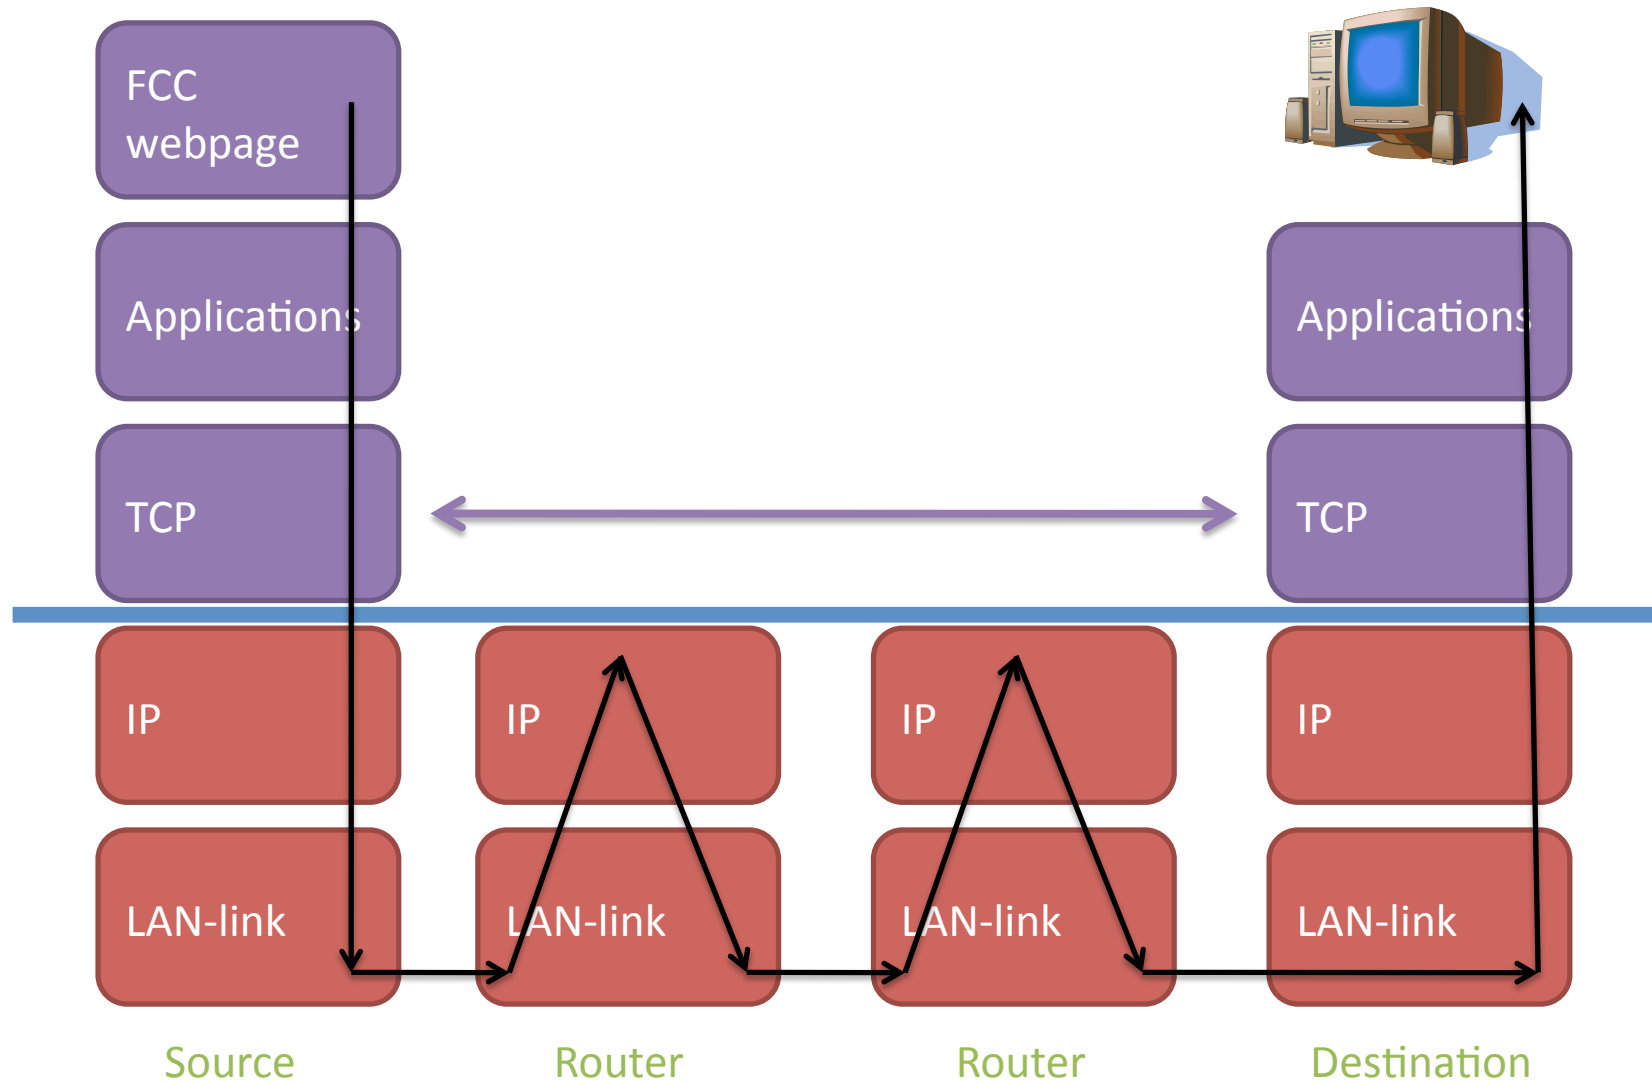
Internet Layering



Peering



TCP: end-to-end



Transport layer problems

- Routers may drop packets
- Packets may arrive out of order
- Destination may not be able to process received packets as fast as they arrive
 - Causing destination buffer overflow
- Routers in between source and destination may experience congestion
 - Routers may queue many packets, resulting in queuing delay
 - Routers may drop packets, resulting in packet loss

Reliable service

- Problem: routers may drop packets
- Solution: transport layer retransmits dropped packets
 - *if* want reliable service at transport layer

Connection-oriented service

- Problem: packets may arrive out of order
- Solution: transport layer puts packets in order before handing them to the application layer
 - *if* want connection-oriented service at transport layer

Flow control

- Problem: destination may not be able to process received packets as fast as they arrive
 - Causing destination buffer overflow
- Solution: slow down source rate to match destination rate
 - Called “flow control”

Congestion control

- Problem: routers in between source and destination may experience congestion
 - Routers may queue many packets, resulting in queuing delay
 - Routers may drop packets, resulting in packet loss
- Solution: slow down rates of sources causing congestion
 - Called “congestion control”

Interface to application layer

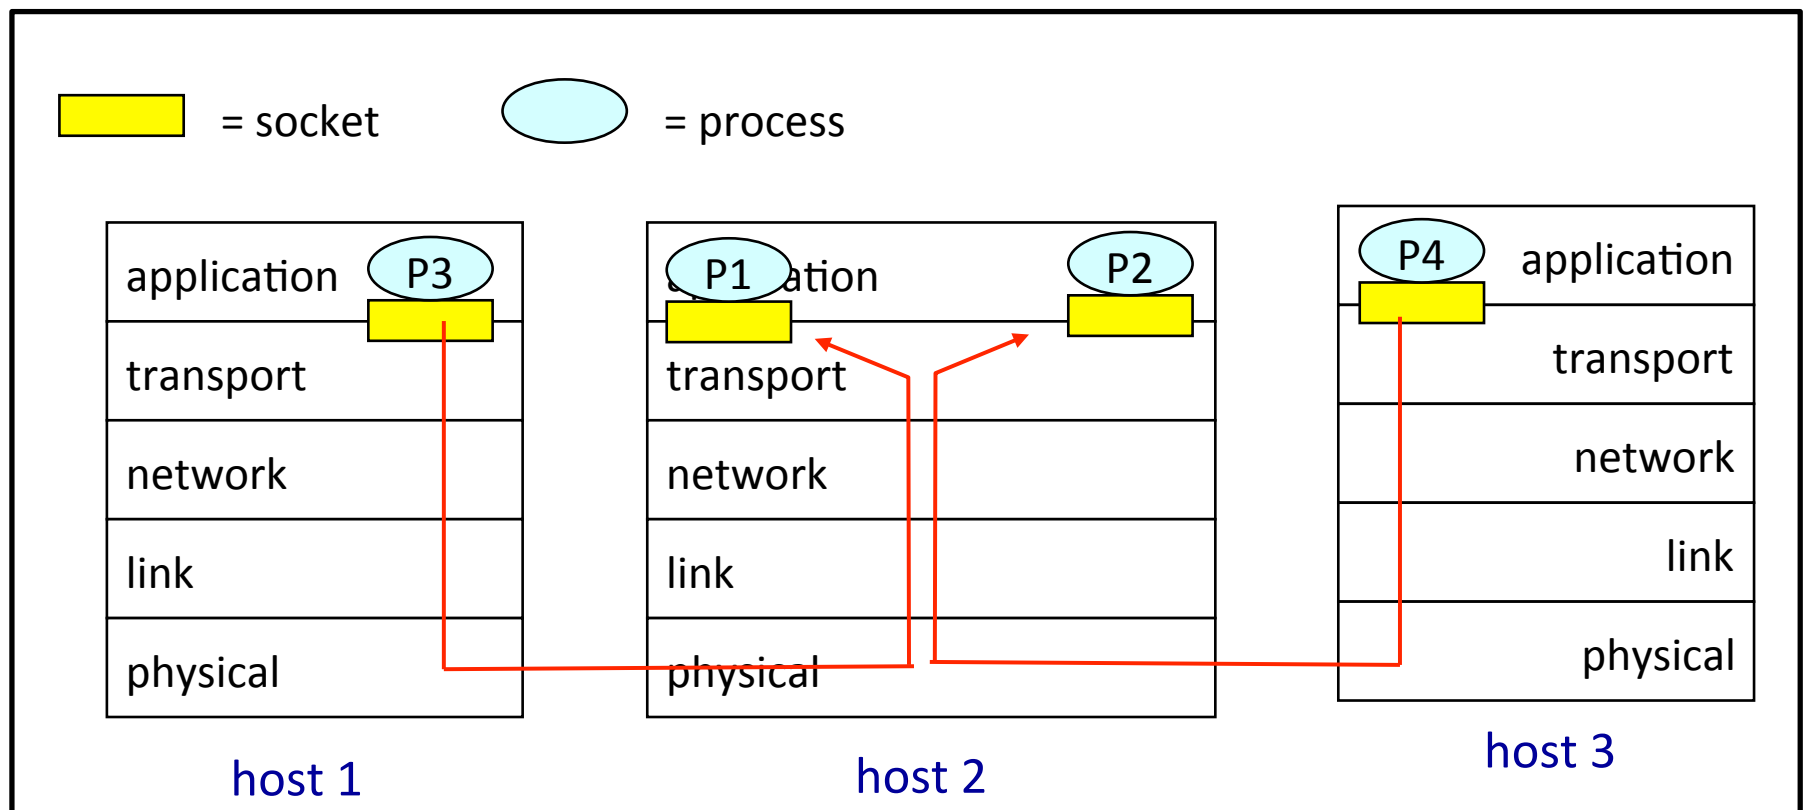
- Source side: accept application-layer messages from various processes
 - Note application choice reliable/unreliable, connection-oriented/connectionless, flow control, congestion control
 - “multiplexing”
- Destination side: feed received segments to various application layer processes
 - “demultiplexing”

Interface: ports

- Part of the interface from application layer to transport layer
- Multiple processes may be talking over Internet at the same time
- Among all packets coming into a machine, separate them by “port”
 - Port roughly corresponds to an application, e.g. port 80 for http, port 25 for smtp, ...

Interface: sockets


- Interface to application layer through a “socket”



Interface: connectionless sockets

- Source side multiplexing
 - Unique application process for each active source port number
 - Application for a source port passes IP address, source port number, destination port number to transport layer
 - Transport layer puts source & destination port numbers in packet header, and passes destination IP address to network layer
- Destination side demultiplexing
 - Obtain destination port number from header
 - Pass to socket = (destination IP address, destination port)
 - to application process associated with destination port

Interface: connectionless sockets

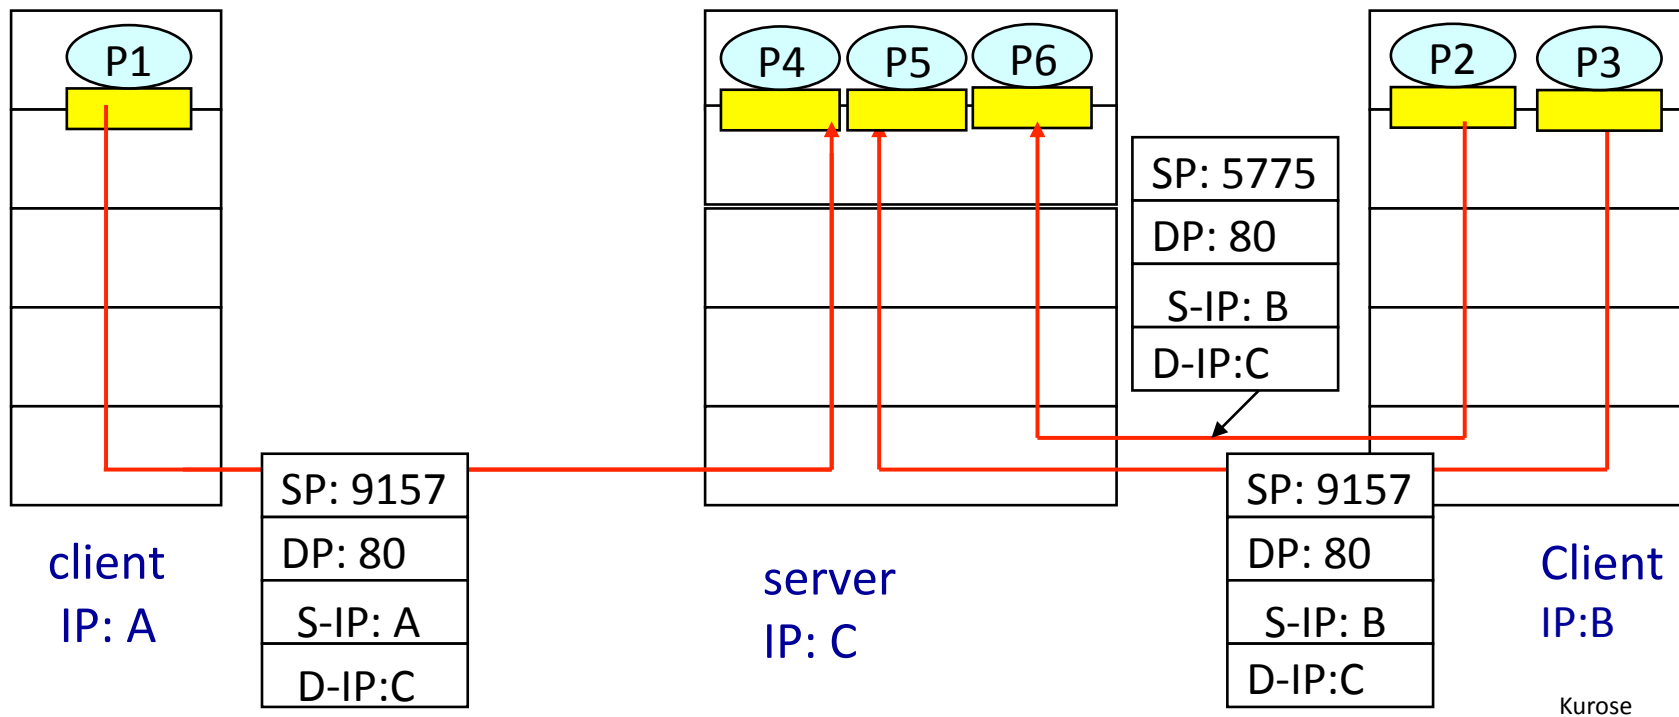


Kurose

Interface: connection-oriented sockets

- Source side multiplexing
 - Unique application process for each active combination of (source port number, destination port number, destination IP address)
 - Application passes IP address, source port number, destination port number to transport layer
 - Transport layer puts source & destination port numbers in packet header, and passes destination IP address to network layer
- Destination side demultiplexing
 - Obtain source and destination port numbers from TCP header
 - Obtain sequence number from TCP header
 - Request dropped packets
 - Put packets in order
 - Pass to socket = (src IP, src port, dest IP, dest port)
 - to application process associated with (src IP, src port, dest port)

Interface: connection-oriented sockets

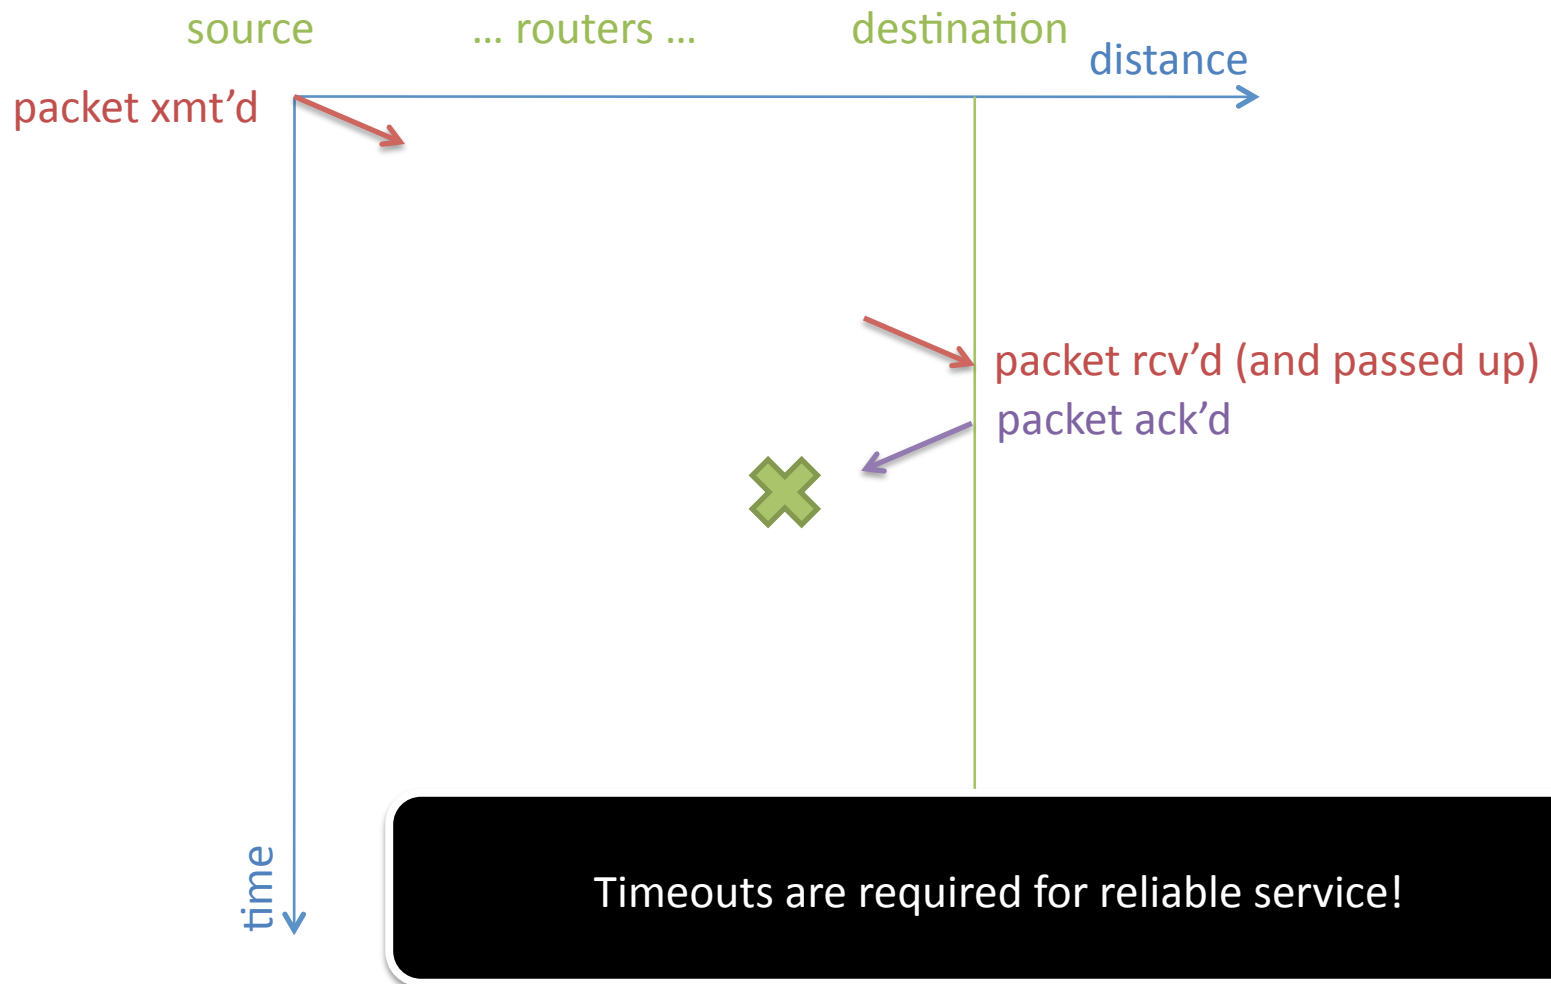


Kurose

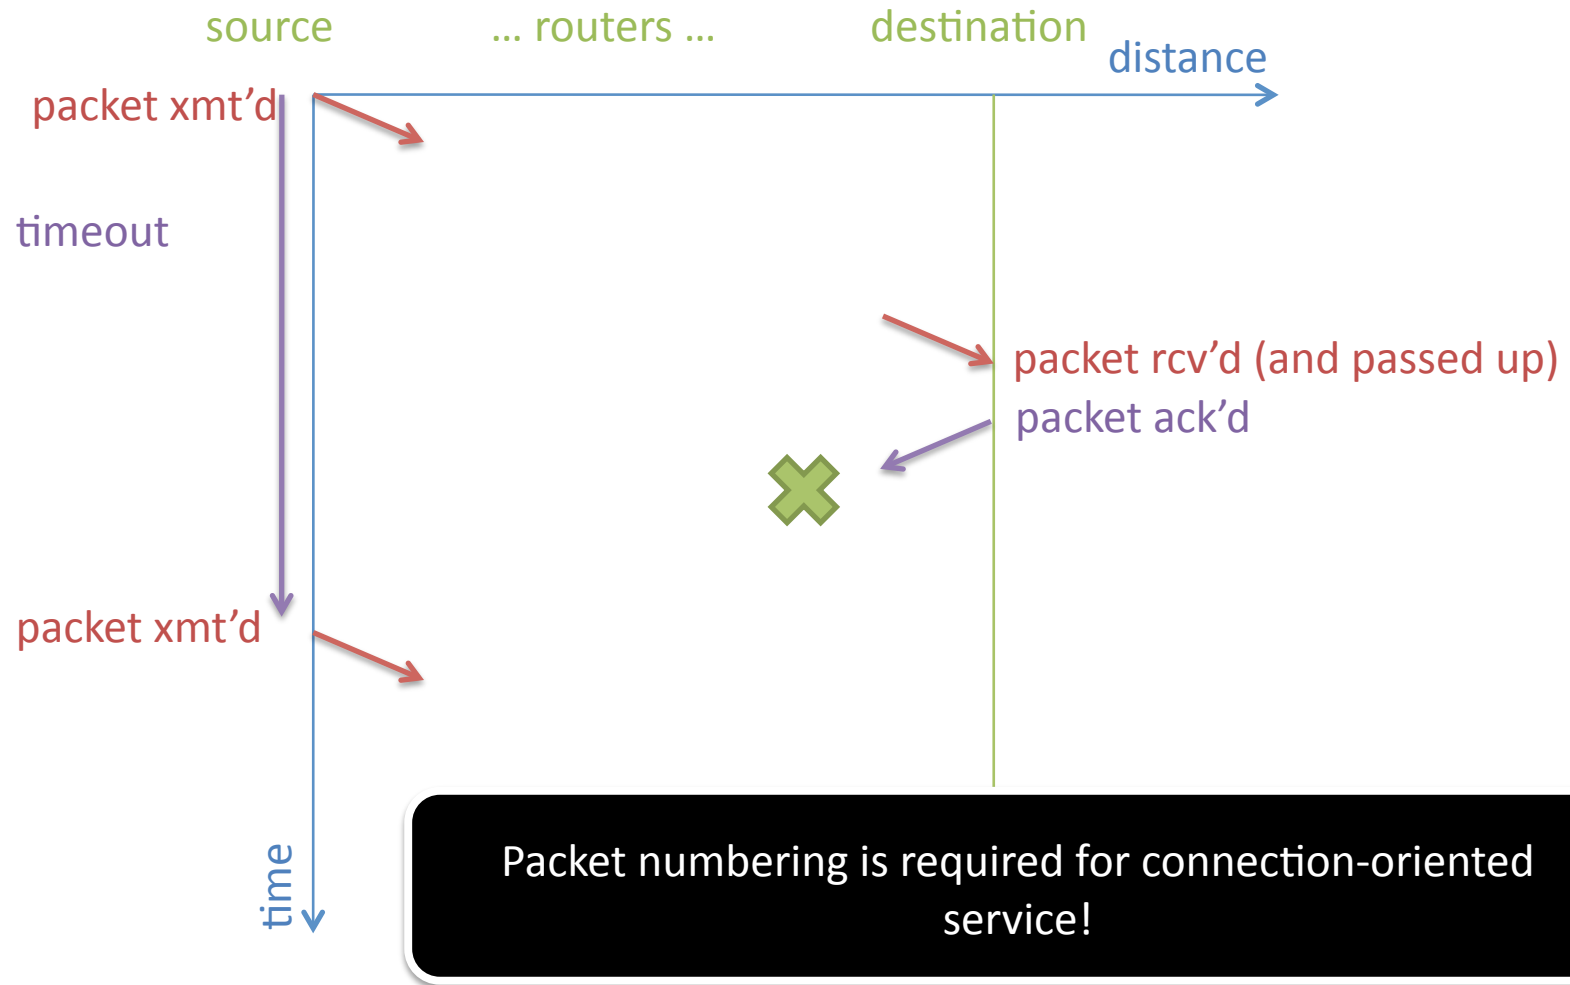
essential elements

- Goals:
 - Reliable
 - Connection-oriented
- Components:
 - Packet numbering
 - Feedback
 - Retransmission

timeouts




packet numbering



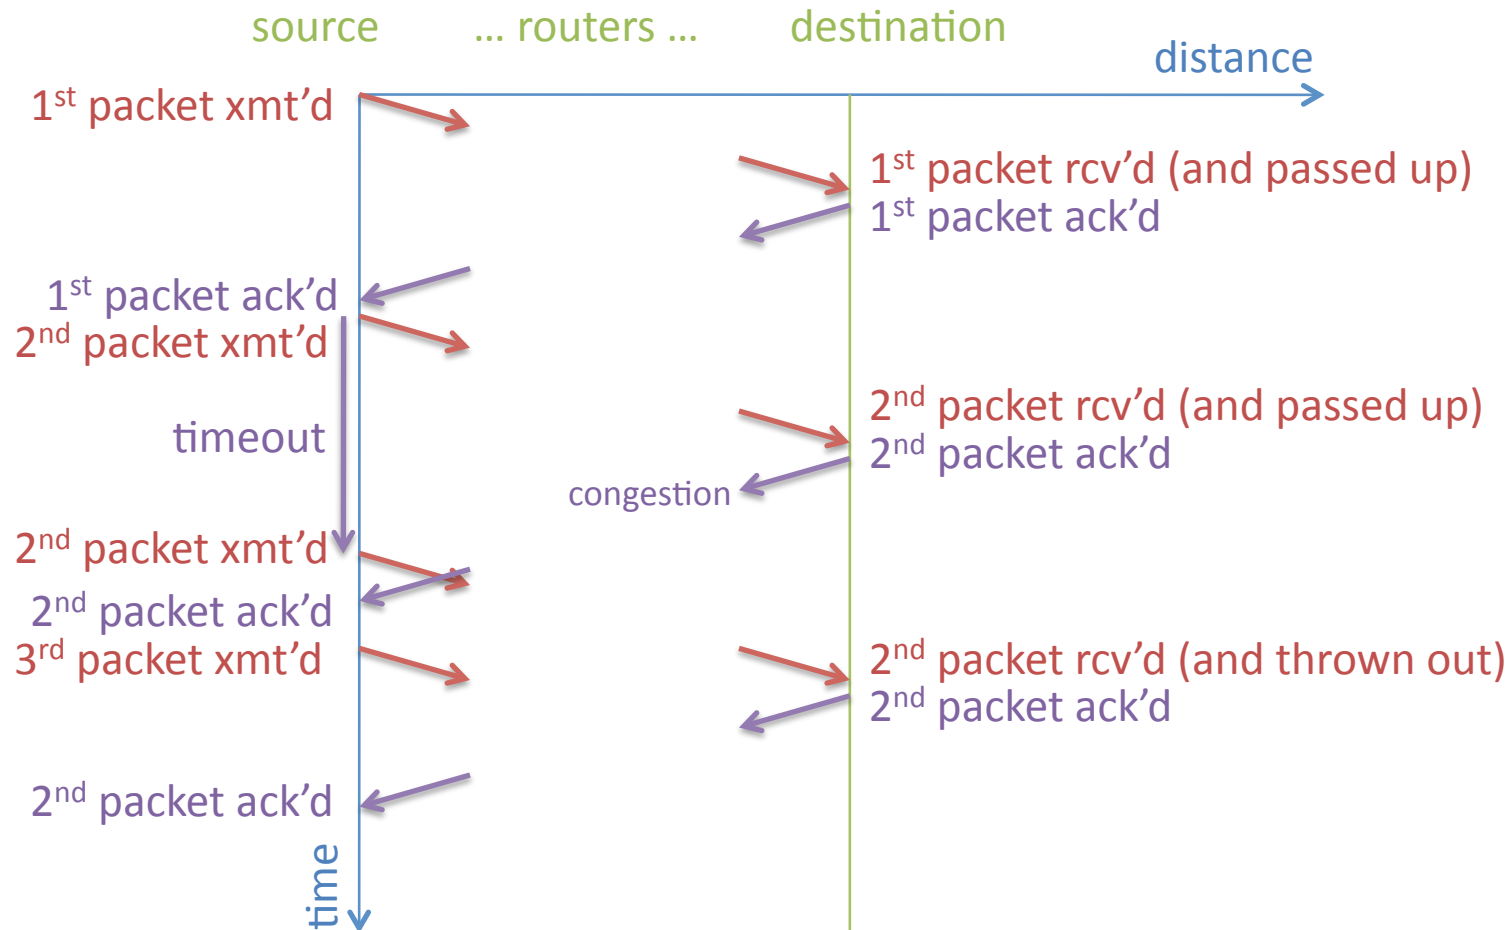
Flow Control Window

- Goals:
 - Reliable
 - Connection-oriented
 - Flow control
- Method:
 - Source allowed to transmit only W packets at a time
 - Must wait for ACKs of previous packets before transmitting more packets (flow control)
 - Packets are numbered to detect loss
 - Packets not ACK'd within *timeout* are retransmitted

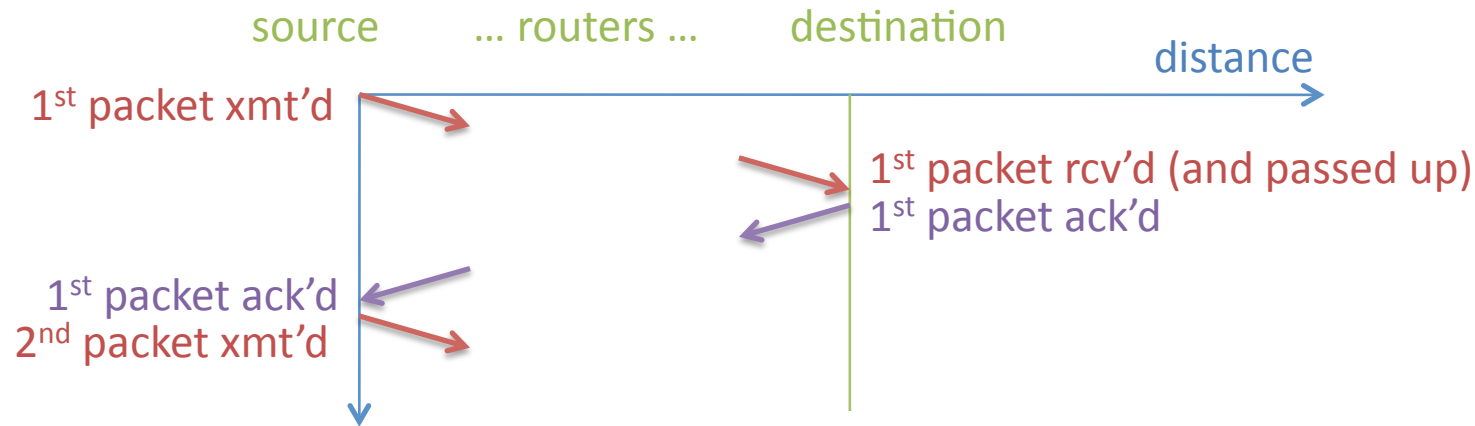
Window flow control



Alternating Bit Protocol ($W=1$)



Efficiency of ABP?



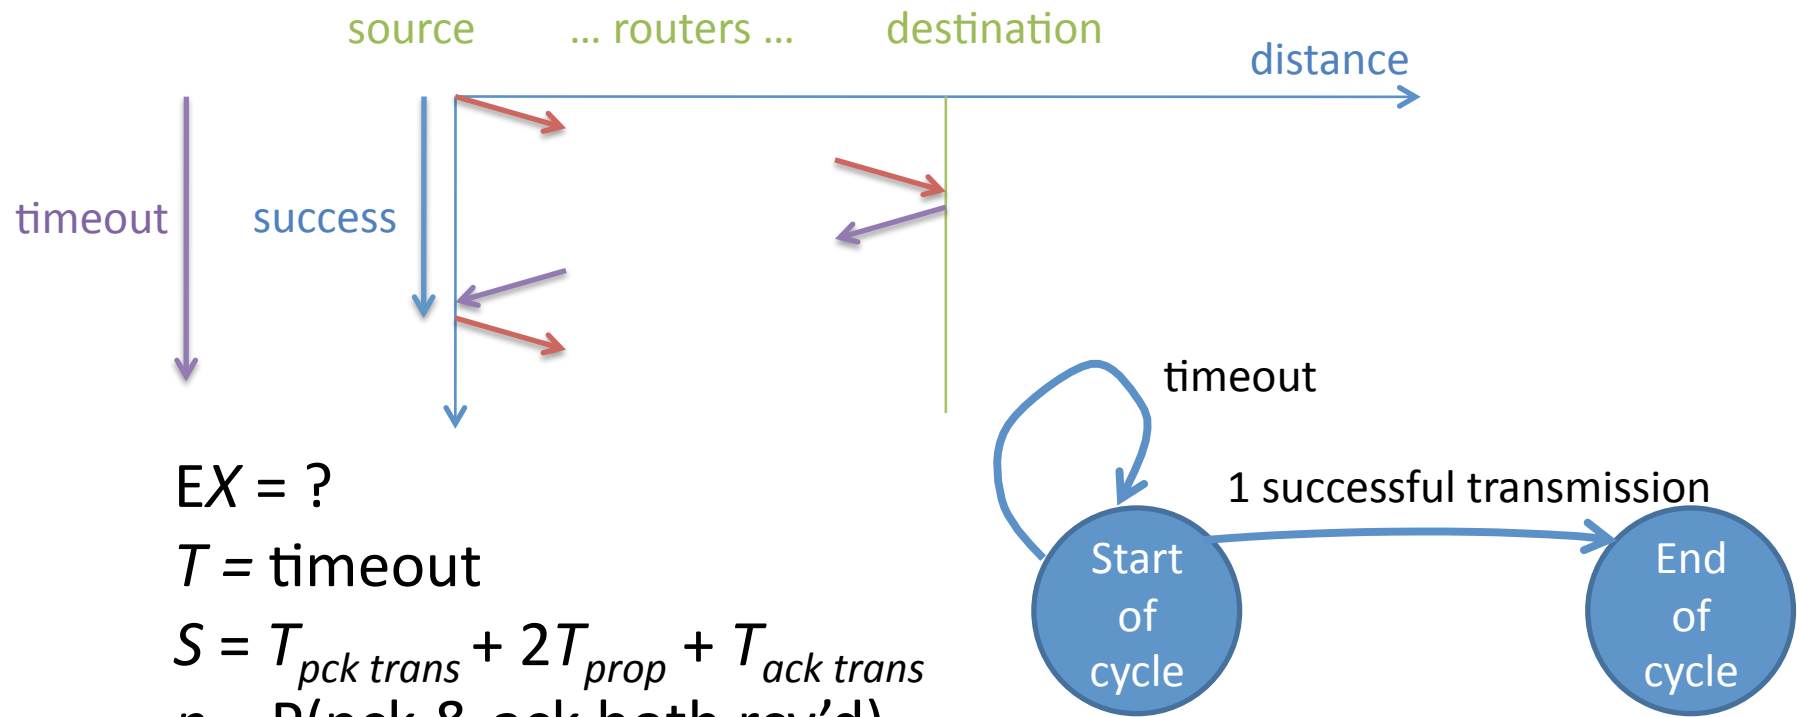
η = proportion of time used for successful transmissions

$$= T_{trans} / (EX)$$

where EX = average time to successfully transmit a packet

$EX = ?$

Efficiency of ABP



$EX = ?$

$T = \text{timeout}$

$S = T_{pck\ trans} + 2T_{prop} + T_{ack\ trans}$

$p = P(\text{pck \& ack both rcv'd})$

$EX = p(?) + (1-p)(?)$

$EX = p(S) + (1-p)(T+EX)$

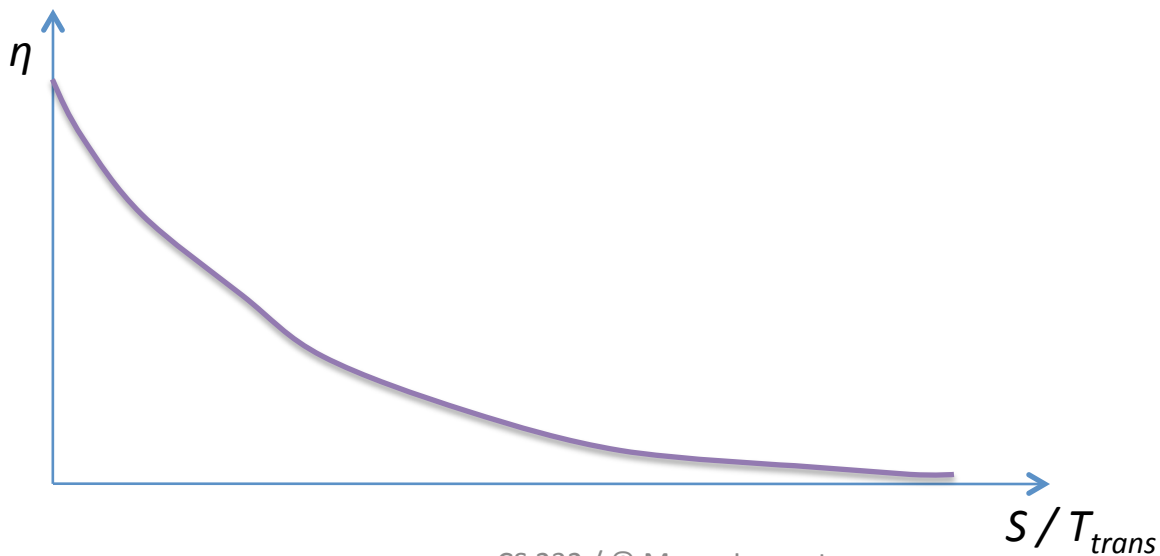
So $EX = S + T(1-p)/p$

Efficiency of ABP


$$EX = S + T(1-p)/p$$

$$\eta = T_{trans} / (EX)$$


$$= T_{trans} / [S + T(1-p)/p]$$



Selective Repeat Protocol ($W > 1$, fixed)



Selective Repeat Protocol (example with $W=4$)



Internet transport layer

Two options at Transport Layer:

- User Datagram Protocol (UDP)
 - No flow control (at transport layer)
 - No congestion control (at transport layer)
 - i.e. no transport layer mechanism to slow down source
- Transmission Control Protocol (TCP)
 - Provides flow control at transport layer
 - Provides congestion control at transport layer
 - i.e. source may be slowed down due to destination speed or due to congestion in routers

UDP flow and congestion control

No flow or congestion control at transport layer

- No UDP flow control:
 - Source *application layer* must implement its own flow control
 - using whatever information that the destination application layer sends to it
- No UDP congestion control
 - Source *application layer* should implement its own congestion control
 - *if* it wants to be nice to the network

UDP packet acks

No transport layer packet acknowledgements

- Destination UDP doesn't tell source UDP if it received a packet ("unreliable")
 - Since source UDP doesn't know when packets are dropped, it doesn't retransmit dropped packets
 - Destination *application layer* responsible for acknowledgements or other feedback to source, if desired
- Destination UDP doesn't put packets back into transmitted order ("connectionless")
 - Destination *application layer* responsible for packet reordering

TCP flow and congestion control

Both flow and congestion control at transport layer

- Flow control:
 - Destination TCP signals when it is ready to accept more packets
- Congestion control
 - Source TCP speeds up or slows down based on estimate of congestion

TCP packet acks

Transport layer issues packet acknowledgements

- Destination TCP tells source TCP if it received a packet (and which packet) using acknowledgements (ACKs) (“reliable”)
 - Source TCP retransmits packets which are not ACK’d within a *timeout*
- Destination TCP reorders packets into their original order (“connection-oriented”)
 - Before handing them to destination application layer

TCP flow and congestion control

Use ACKs from destination TCP to:

- Know when destination is ready to accept more packets (flow control)
 - When the source receives an ACK, it knows that destination has processed that packet
 - Specific field in reverse direction packets' headers
- Estimate congestion (congestion control)
 - The rate at which the source receives ACKs tells it something about the delay between the source and destination

Congestion control

- Problem: routers in between source and destination may experience congestion
 - Routers may queue many packets, resulting in queuing delay
 - Routers may drop packets, resulting in packet loss
- Solution: slow down rates of sources causing congestion
 - Called “congestion control”

Window-based congestion control


- Goals:
 - Reliable
 - Connection-oriented
 - Flow control
 - Congestion control
- Method:
 - Source allowed to transmit only W packets at a time
 - Must wait for ACKs of previous packets before transmitting more packets (flow control)
 - Packets are numbered to detect loss
 - Packets not ACK'd within *timeout* are retransmitted
 - Window size W is modified based on estimate of congestion (congestion control)

TCP-Tahoe

Algorithm for W :

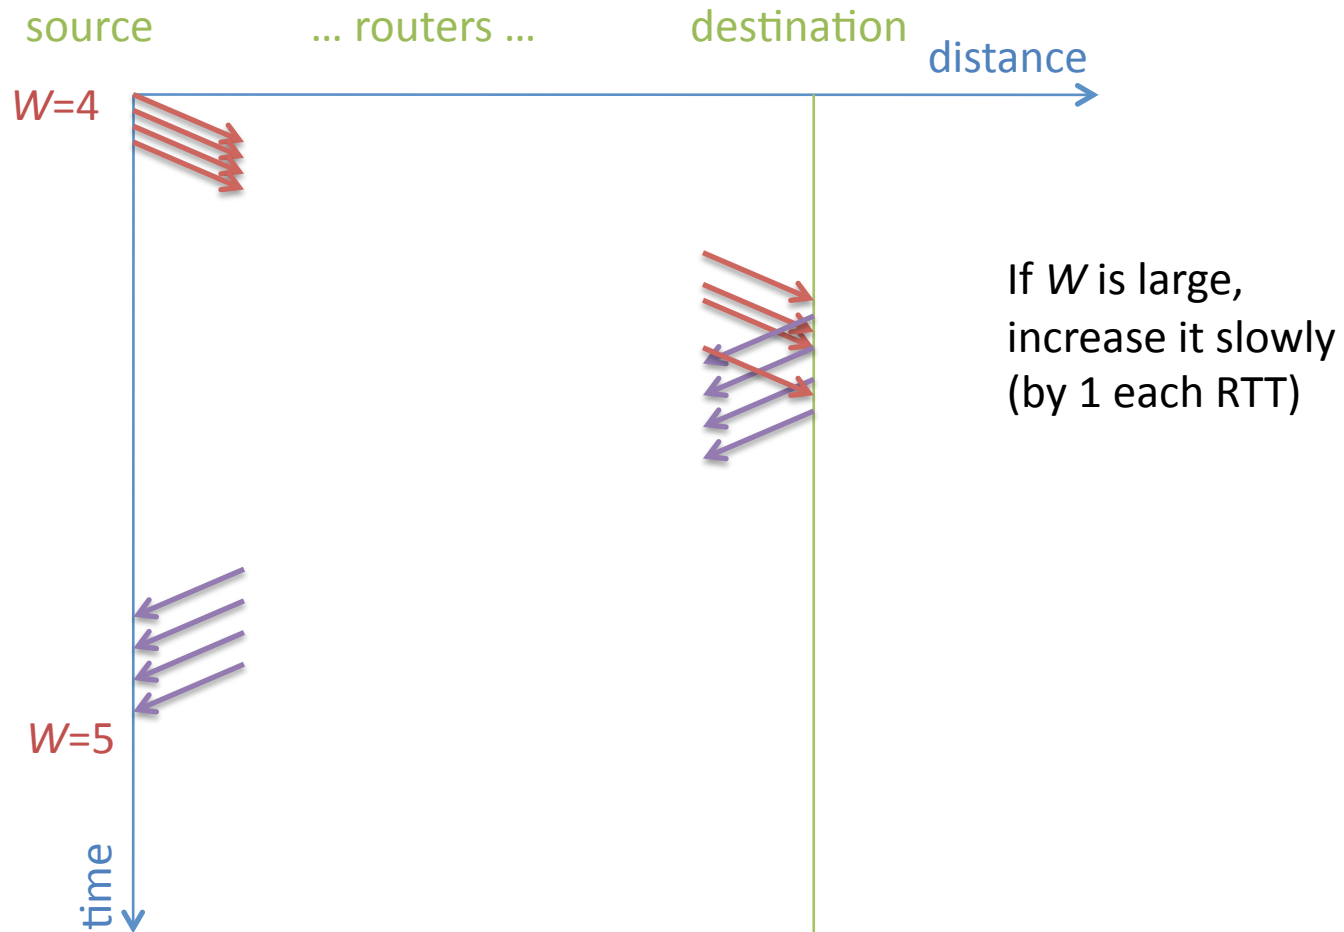
- Start timidly with $W=1$.
- Increase W when no congestion
 - “Slow Start”: If W is small, increase it quickly (by 1 for each received ACK)
 - “Congestion Avoidance”: If W is large, increase it slowly (by 1 each RTT)
 - There is an upper limit on W
- Decrease W when congestion
 - If there is a timeout, go back to $W=1$

Increasing W when it is small (slow start)




Increasing W when it is large


Congestion control




TCP congestion control




Cumulative ACK



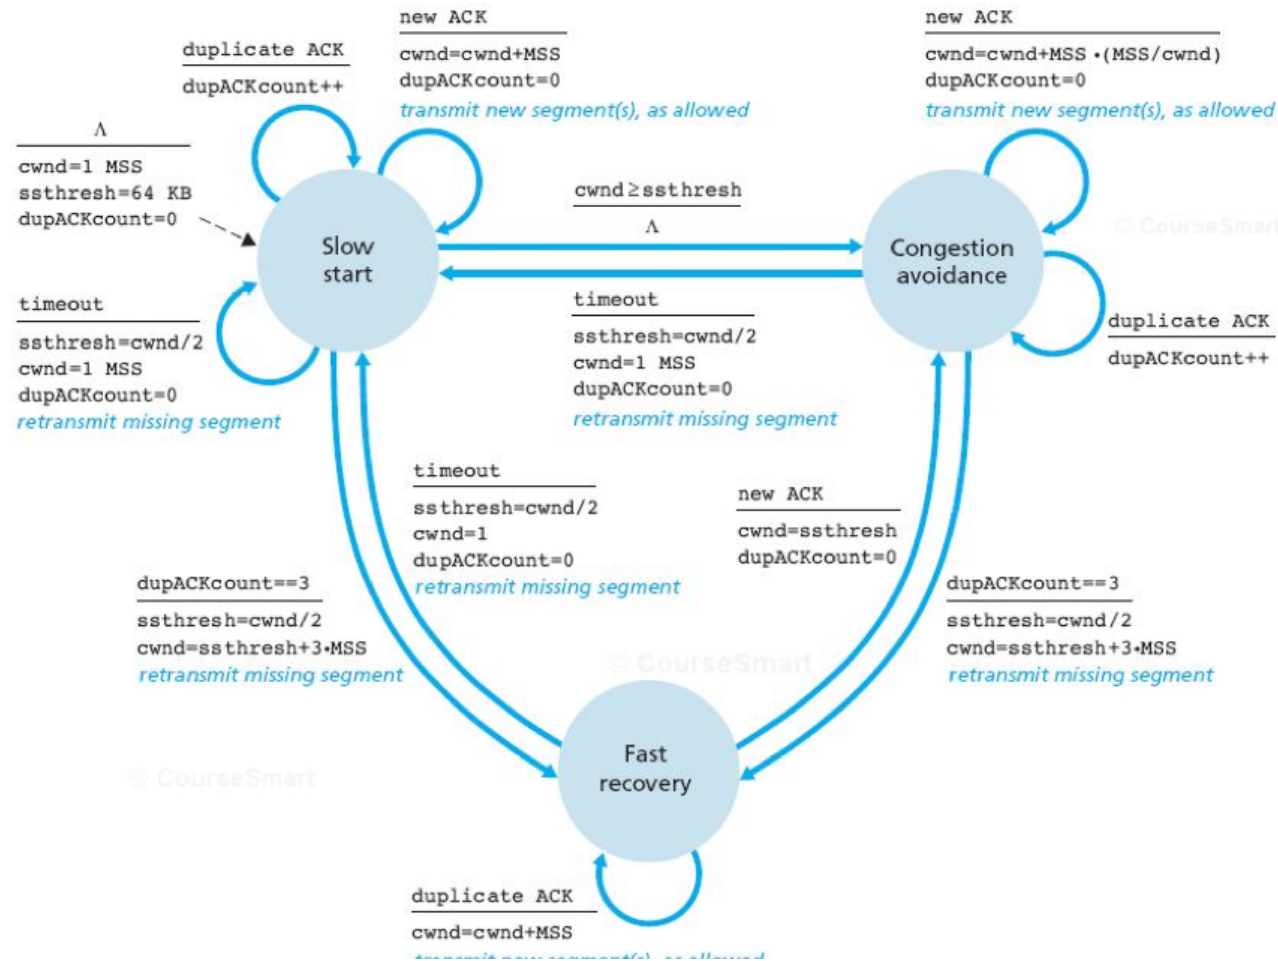
TCP: fast retransmit



TCP-Reno Fast Recovery



Finite State Machine



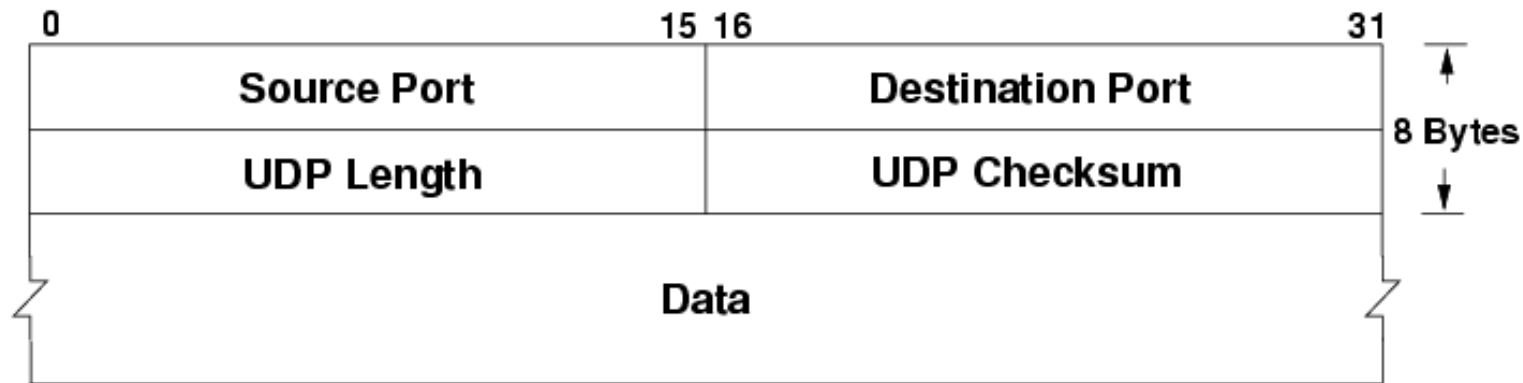
TCP: additional details

- Actual version:
 - uses “segments” and byte #s instead of pkt #s
 - keeps a timer only for the oldest segment
 - dynamically adjusts the threshold between slow start and congestion avoidance
 - dynamically adjusts timeout value
 - has additional rules regarding what to ACK when
 - “TCP-NewReno” has additional tweaks
 - there are many other variants ...

UDP vs. TCP

- TCP used by:
 - Applications that want packet retransmission & packet reordering
 - And are willing to put up with congestion control
 - e.g. email, http, ftp
- UDP used by:
 - Applications that are not willing to put up with congestion control
 - And applications that don't need packet retransmission & packet reordering
 - e.g. streaming, voice over IP (VoIP), video conferencing, gaming


UDP header



Wikipedia Commons

- ports
- src & dest IP addresses in IP header
- no packet numbers

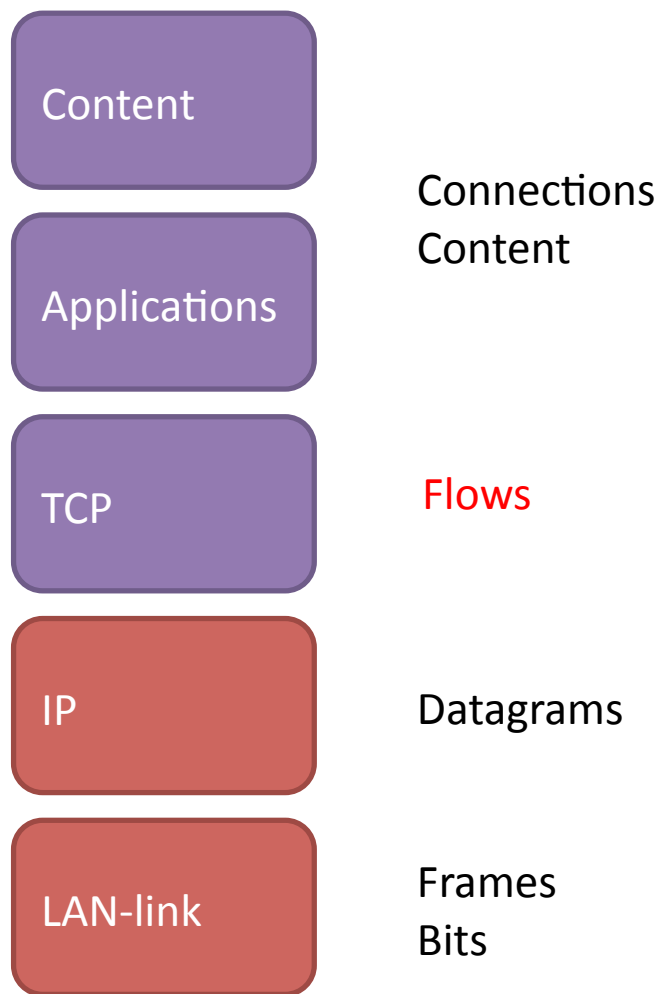
TCP header



- ports
- sequence & ACK numbers for retransmission & ordering
- window size for flow & congestion control

Wikipedia Commons

Traffic characterization by layer



Transport layer models

- Flow models
 - Multiple flows
 - Graph of
 - end points
 - links & nodes:
 - all links & nodes
 - abstracted set of links & nodes
 - or just bottleneck links / nodes

Transport layer models

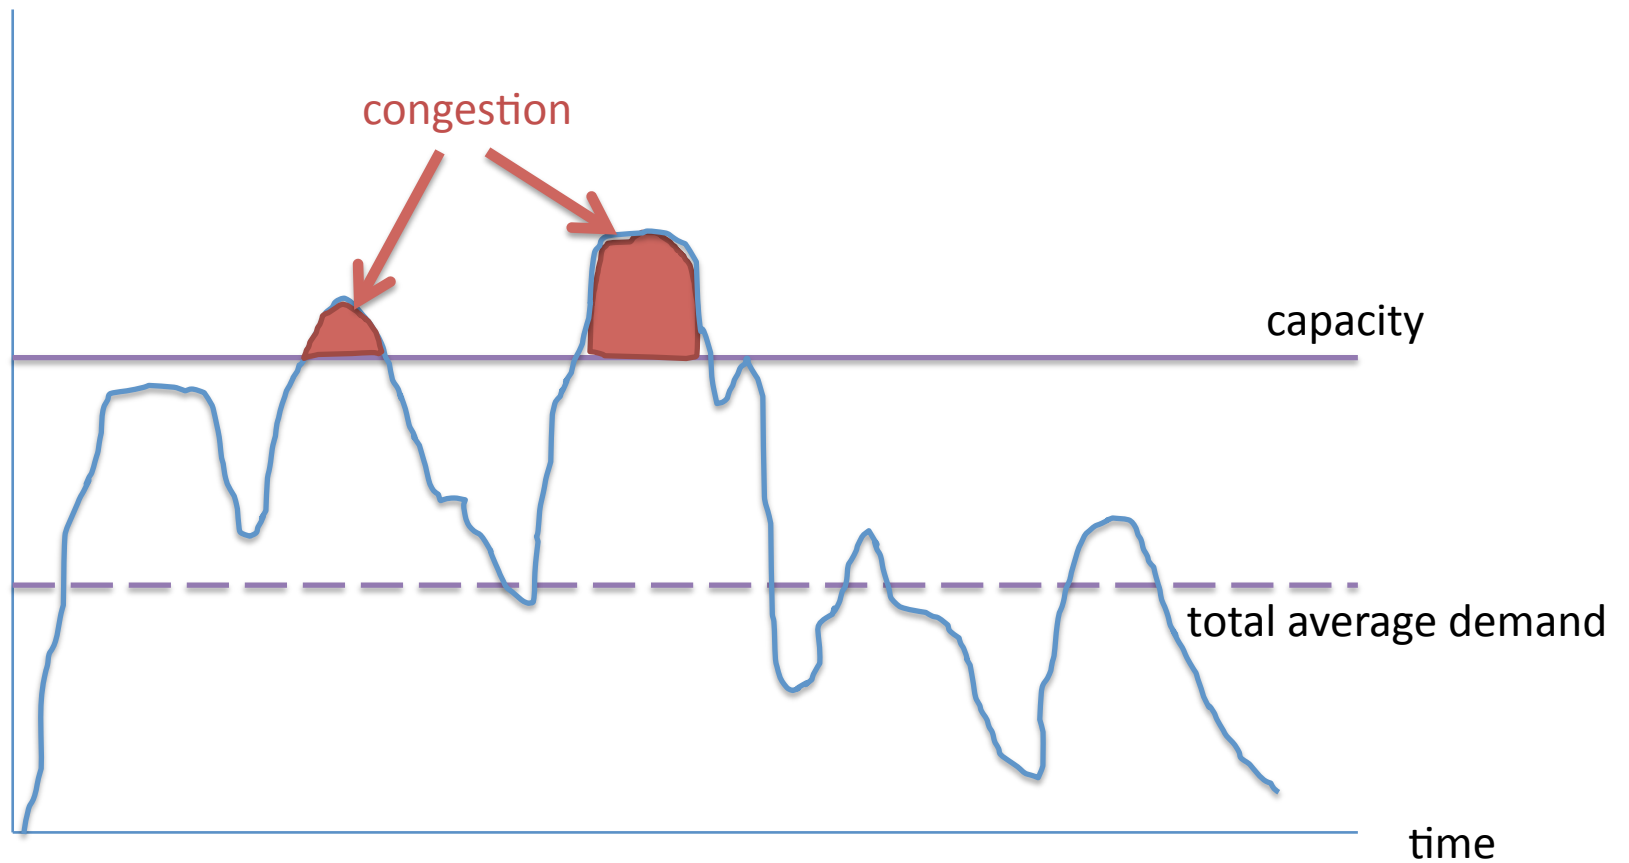
- Flow initiation & termination
 - by application layer traffic model
 - or considered under a fixed number of connections

Transport layer models

- Flow rate
 - Dynamic capacity of links & nodes given by lower layer models
 - Or may be abstracted by an independent model, e.g. flow restrictions & delay
 - Plus model of window or rate flow & congestion control


Congestion

demand in bits per second




Multiple flow model

- What is the throughput on a single connection?
 - $\text{Min}(R_s, R/3, R_c)$



Kurose fig. 1.20

Multiple flow model

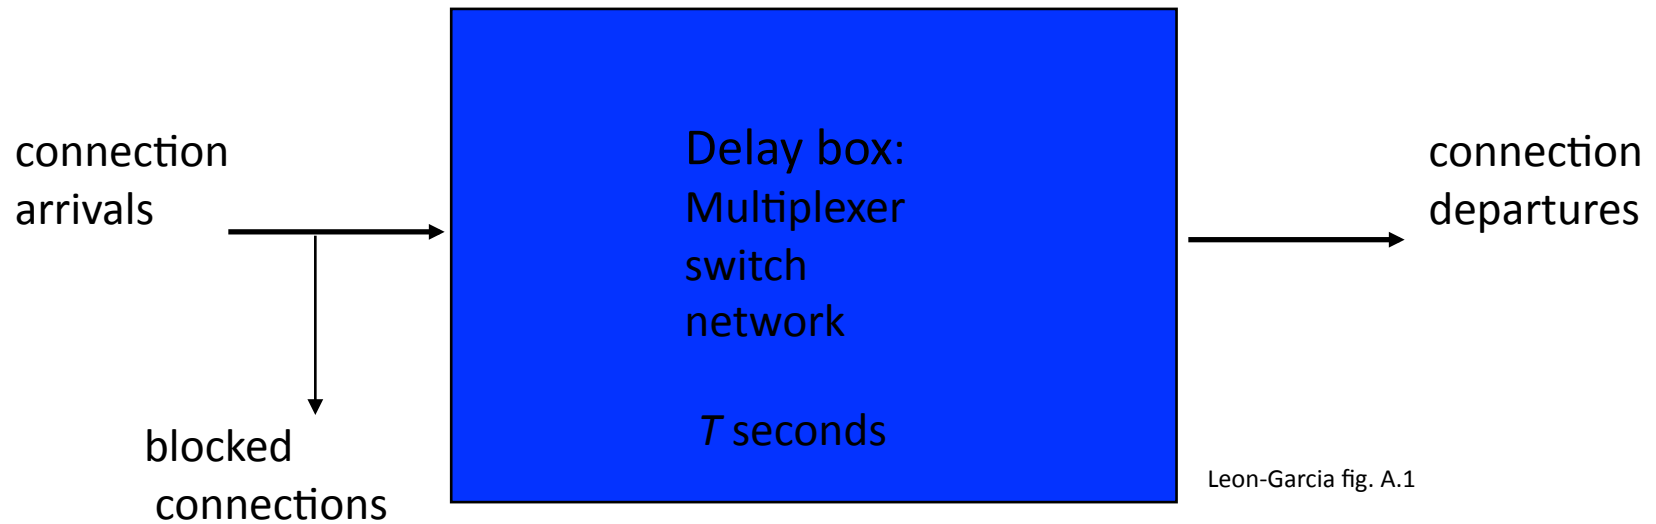


Kurose fig. 3.56

Transport layer model examples


- TCP flow model
 - TCP protocol
 - Model of pertinent events, e.g. timeouts
 - either from lower layer traffic models
 - or abstracted
 - e.g. for http, email, some streaming
- UDP flow model
 - Model of pertinent events, e.g. packet losses
 - either from lower layer traffic models
 - or abstracted
 - e.g. for VoIP, some streaming

Connection models



- Arrivals
- Duration
- Connection access control

Connection Queuing Model



Leon-Garcia fig. A.6