


# CS273A: Ensemble Methods



Prof. Alexander Ihler  
Fall 2023

# Ensemble methods

- Why learn one classifier when you can learn many?
- Ensemble: combine many predictors
  - (Weighted) combinations of predictors
  - May be same type of learner or different



**Various options for getting help:**



**“Who wants to be a millionaire?”**

# Ensemble Methods

---

## Basic Ensembles

Committees

Stacking

Mix of Experts

Bagging

Gradient Boosting

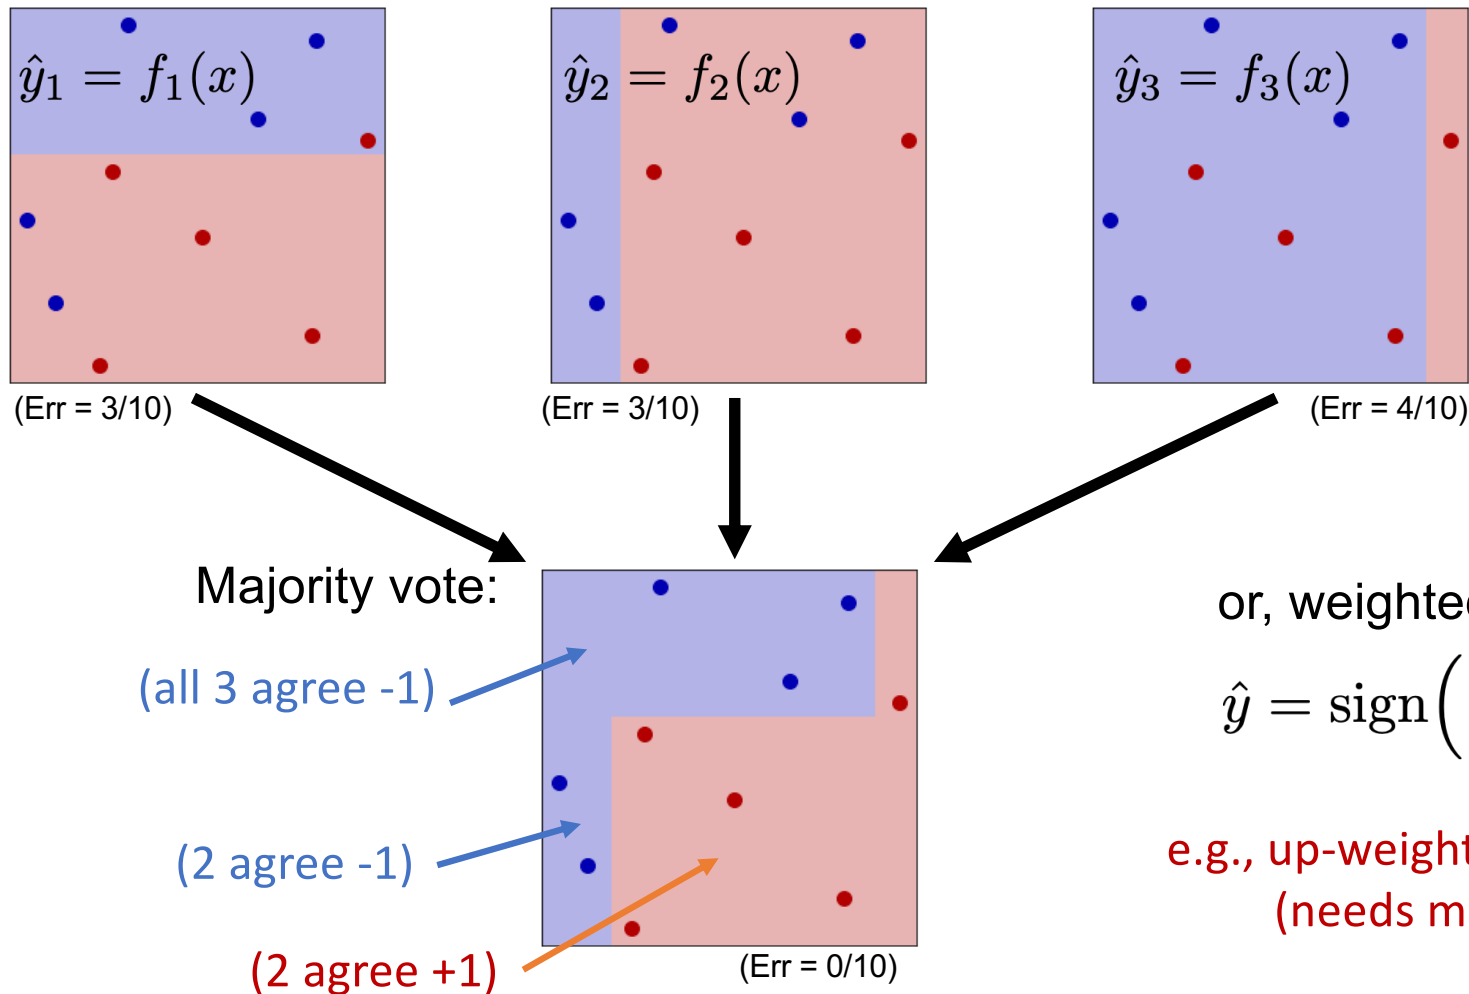
AdaBoost

# Simple ensembles

- “Committees”

- Average / majority vote of several predictors

$$y \in \{-1, +1\}$$



# “Stacked” ensembles

---

- Train a “predictor of predictors”
  - Treat individual predictors as features

$$\hat{y}_1 = f_1(x_1, x_2, \dots)$$

$$\hat{y}_2 = f_2(x_1, x_2, \dots) \quad \Rightarrow \quad \hat{y}_e = f_e(\hat{y}_1, \hat{y}_2, \dots)$$


...

- Similar to multi-layer perceptron idea
- Special case: binary,  $f_e$  linear  $\Rightarrow$  weighted vote
- Can train stacked learner  $f_e$  on validation data
  - Avoids giving high weight to overfit models

# Mixtures of experts

- Can make weights depend on  $x$ 
  - Weight  $\alpha_z(x)$  indicates “expertise”
  - Combine using weighted average (or even just pick largest)

Example:



Mixture of three linear predictor experts

Weighted average:

$$f(x; \omega, \theta) = \sum_z \alpha_z(x; \omega) f_z(x; \theta_z)$$

Weights: (multi) logistic regression

$$\alpha_z(x; \omega) = \frac{\exp(x \cdot \omega^z)}{\sum_c \exp(x \cdot \omega^c)}$$

If loss, learners, weights are all differentiable, can train jointly...

# Ensemble Methods

---

## Basic Ensembles

Committees

Stacking

Mix of Experts

## Bagging

## Gradient Boosting

## AdaBoost

# Ensemble methods

---

- Where can we get a diverse collection of learners?
  - Maybe create one artificially?
- “Bagging” = bootstrap aggregation
  - Learn many classifiers, each with only part of the data
  - Combine through model averaging
- Remember overfitting: “memorize” the data
  - Used test data to see if we had gone too far
  - Cross-validation
    - Make many splits of the data for train & test
    - Each of these defines a classifier
    - Typically, we use these to check for overfitting
    - Could we instead combine them to produce a better classifier?




# Bagging

---

- Bootstrap
  - Create a random subset of data by sampling
  - Draw  $m'$  of the  $m$  samples, with replacement (some variants w/o)
    - Some data left out; some data repeated several times
- Bagging
  - Repeat  $K$  times
    - Create a training set of  $m' \leq m$  examples
    - Train a classifier on the random training set
  - To test, run each trained classifier
    - Each classifier votes on the output, take majority
    - For regression: each regressor predicts, take average
- Notes:
  - Some complexity control: harder for each to memorize data
  - Doesn't work for linear models (average of linear functions is linear function), but perceptrons OK (linear + threshold = nonlinear)


# Bias / variance

“The world”      Data we observe



$$\hat{y}(x) = \hat{\theta}_0 + \hat{\theta}_1 x$$







- We only see a little bit of data
- Can decompose error into two parts
  - Bias – error due to model choice
    - Can our model represent the true best predictor?
    - Gets better with more complexity
  - Variance – randomness due to data size
    - Better w/ more data, worse w/ complexity



# Bagged decision trees


- Randomly resample data
- Learn a decision tree for each
  - No max depth = very flexible class of functions
  - Learner is low bias, but high variance

Sampling:  
simulates “equally likely”  
data sets we could have  
observed instead, &  
their classifiers




# Bagged decision trees


- Average over collection
  - Classification: majority vote
- Reduces memorization effect
  - Not every predictor sees each data point
  - Lowers effective “complexity” of the overall average
  - Usually, better generalization performance
  - Intuition: reduces variance while keeping bias low




Avg of 5 trees



Avg of 25 trees



Avg of 100 trees



# Bagging in Python

---

```
# Load data set X, Y for training the ensemble...
```

```
m,n = X.shape
```

```
classifiers = [ None ] * num_bags           # Allocate space for learners
```

```
for b in range(num_bags):
```

```
    # Bootstrap sample a dataset of size "mBag"; typically just pick mBag = m
```

```
    ind = np.floor( mBag * np.random.rand( m ) ).astype(int)
```

```
    Xb, Yb = X[ind,:], Y[ind]                # select the data at those indices
```

```
    classifiers[i] = MyClassifier(Xb, Yb)      # Train a model on data Xi, Yi
```

```
# test on data Xtest
```

```
mTest = Xtest.shape[0]
```

```
predict = np.zeros( (mTest, num_bags) ).    # Allocate space for predictions from each model
```

```
for i in range(num_bags):
```

```
    predict[:,i] = classifiers[i].predict(Xtest)    # Apply each classifier
```

```
# Make overall prediction by majority vote
```

```
predict = np.mean(predict, axis=1) > 0    # if +1 vs -1
```

# Random forests

---

- Bagging applied to decision trees
- Problem
  - With lots of data, we usually learn the same classifier
  - Averaging over these doesn't help!
- Introduce extra variation in learner
  - At each step of training, only allow a (random) subset of features
  - Enforces diversity (“best” feature not available)
  - Keeps bias low (every feature available eventually)
  - Average over these learners (majority vote)

```
# in FindBestSplit(X,Y):  
    for each of a subset of features  
        for each possible split  
            Score the split (e.g. information gain)  
        Pick the feature & split with the best score  
    Recurse on left & right splits
```

# Ensemble Methods

---

## Basic Ensembles

Committees

Stacking

Mix of Experts

## Bagging

## Gradient Boosting

## AdaBoost

# Ensembles

---

- Weighted combinations of predictors
- “Committee” decisions
  - Trivial example
  - Equal weights (majority vote / unweighted average)
  - Might want to weight unevenly – up-weight better predictors
- Boosting
  - Focus new learners on examples that others get wrong
  - Train learners sequentially
  - Errors of early predictions indicate the “hard” examples
  - Focus later predictions on getting these examples right
  - Combine the whole set in the end
  - Convert many “weak” learners into a complex predictor




# Gradient boosting

- Learn a regression predictor
- Compute the error residual
- Learn to predict the residual


Learn a simple predictor...

$$f_1(x^{(i)}) \approx y^{(i)}$$



Then try to correct its errors

$$\epsilon^{(i)} = y^{(i)} - f_1(x^{(i)})$$



# Gradient boosting

- Learn a regression predictor
- Compute the error residual
- Learn to predict the residual


$$f_1(x^{(i)}) \approx y^{(i)}$$

$$\epsilon^{(i)} = y^{(i)} - f_1(x^{(i)})$$

$$f_2(x^{(i)}) \approx \epsilon^{(i)}$$


Combining gives a better predictor...

$$\Rightarrow f_1(x^{(i)}) + f_2(x^{(i)}) \approx y^{(i)}$$



Can try to correct its errors also, & repeat

$$\epsilon_2^{(i)} = y^{(i)} - f_1(x^{(i)}) - f_2(x^{(i)}) \quad \dots$$




# Gradient boosting


- Learn sequence of predictors
- Sum of predictions is increasingly accurate
- Predictive function is increasingly complex

$$y^{(i)} \approx \sum_z f_z(x^{(i)})$$

Data & prediction function



Error residual



...

# Gradient boosting

---

- Make a set of predictions  $\hat{y}[i]$
- The “error” in our predictions is  $J(y, \hat{y})$
- For MSE:  $J(.) = \sum (y[i] - \hat{y}[i])^2$
- We can “adjust”  $\hat{y}$  to try to reduce the error
- $\hat{y}[i] = \hat{y}[i] + \alpha f[i]$
- $f[i] \approx \nabla J(y, \hat{y}) = (y[i] - \hat{y}[i])$  for MSE
- Each learner is estimating the gradient of the loss function
- Gradient descent: take sequence of steps to reduce  $J$
- Sum of predictors, weighted by step size  $\alpha$

# Gradient boosting (classification)


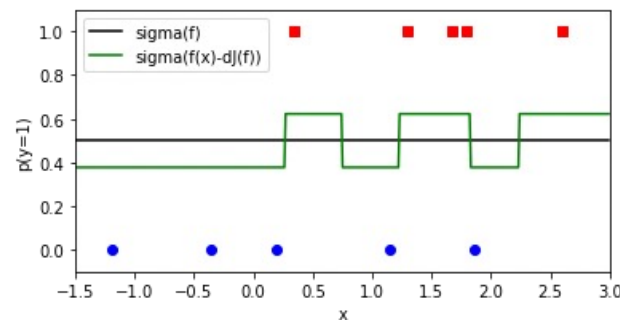

- Ex: “logistic regression” using boosting
  - “Response”  $r(x)$  (standard LR: a linear f’n of  $x$ )
  - Probability  $\sigma(r) \in [0,1]$
  - Loss  $J(r) = \sum_i y^{(i)} \log r(x^{(i)}) + (1-y^{(i)}) \log (1-r(x^{(i)}))$

Learn a simple predictor... Find the loss gradient:

$$r_0(x^{(i)}) = 0.5$$

$$\frac{\partial J}{\partial r^{(i)}} = y^{(i)} - \sigma(r^{(i)})$$

Learn to approximate it:



# Gradient boosting in Python

```
# Load data set X, Y ...
```

```
learner = [None] * num_boost      # storage for ensemble of models
```

```
alpha = [1.0] * num_boost        # and weights of each learner
```

```
mu = Y.mean()                    # often start with constant "mean" predictor
```

```
dJ = Y - mu                      # subtract this prediction away (assumes MSE)
```

```
for b in range( num_boost ):
```

```
    learner[b] = MyRegressor( X, dJ )    # regress to predict gradient direction dJ using X
```

```
    alpha[b] = 1.0                  # alpha: "learning rate" or "step size"
```

```
    # smaller alphas need to use more classifiers, but may predict better given enough of them
```

```
    # compute the residual given our new prediction:
```

```
    dJ = dJ - alpha[b] * learner[b].predict(X).    # update gradient (assumes MSE loss)
```

```
# test on data Xtest
```

```
mTest = Xtest.shape[0]
```

```
predict = np.zeros( (mTest,) ) + mu    # Allocate space for predictions & add 1st (mean)
```

```
for b in range(num_boost):
```

```
    predict += alpha[b] * learner[b].predict(Xtest) # Apply predictor of next residual & accum
```

# Ensemble Methods

---

## Basic Ensembles

Committees

Stacking


Mix of Experts

## Bagging


## Gradient Boosting

## AdaBoost


# Boosting example

Original data set,  $D_1$ 


Trained classifier




$$Jw = 0.3, \alpha = 0.42$$

Update weights,  $D_2$ 


Trained classifier



$$Jw = 0.21, \alpha = 0.65$$

Update weights,  $D_3$ 

Trained classifier



$$Jw = 0.18, \alpha = 0.75$$



# Minimizing weighted error

---

- So far we've mostly minimized unweighted error
- Minimizing weighted error is no harder:

Unweighted average loss:

$$J(\theta) = \frac{1}{m} \sum_i J_i(\theta, x^{(i)})$$

Weighted average loss:

$$J(\theta) = \sum_i w_i J_i(\theta, x^{(i)})$$

For any loss (logistic MSE, hinge, ...)

$$J(\theta, x^{(i)}) = (\sigma(\theta x^{(i)}) - y^{(i)})^2$$

$$J(\theta, x^{(i)}) = \max[0, 1 - y^{(i)} \theta x^{(i)}]$$

To learn decision trees, find splits to optimize *weighted* impurity scores:


$p(+1)$  = total weight of data with class +1

$p(-1)$  = total weight of data with class -1  $\Rightarrow H(p)$  = impurity

# Boosting example

Weight each classifier and combine them:

$$.42 * \begin{array}{|c|} \hline \text{blue} \\ \hline \text{red} \\ \hline \end{array} + .65 * \begin{array}{|c|} \hline \text{blue} \\ \hline \text{red} \\ \hline \end{array} + .75 * \begin{array}{|c|} \hline \text{blue} \\ \hline \text{red} \\ \hline \end{array} \gtrless 0$$



1-node decision trees  
“decision stumps”  
*very simple classifiers*

# AdaBoost = “adaptive boosting”

- Pseudocode for AdaBoost

```
# Load data set X, Y ... ; Y assumed +1 / -1
for b in range(num_boost):
    learner[b] = MyClassifier( X, Y, weights=wts ) # train a weighted classifier
    Yhat = learner[b].predict(X)
    e = wts.dot( Y != Yhat ) # compute weighted error rate
    alpha[b] = 0.5 * np.log( (1-e)/e )
    wts *= np.exp( -alpha[b] * Y * Yhat ) # update weights
    wts /= wts.sum() # and normalize them

# Final classifier:
predict = np.zeros( (mTest,) )
for b in range(num_boost):
    predict += alpha[b] * learner[b].predict(Xtest) # compute contribution of each model
predict = np.sign(predict) # and convert to +1 / -1 decision
```

- Notes
  - $e > .5$  means classifier is not better than random guessing
  - $Y * Yhat > 0$  if  $Y == Yhat$ , and weights decrease
  - Otherwise, they increase

# AdaBoost theory

---


- Minimizing classification error was difficult
  - For logistic regression, we minimized MSE or NLL instead
  - Idea: low MSE => low classification error

- Example of a surrogate loss function

- AdaBoost also corresponds to a surrogate loss function


$$C_{ada} = \sum_i \exp[-y^{(i)} f(x^i)]$$

- Prediction is  $\hat{y} = \text{sign}(f(x))$ 
  - If same as  $y$ , loss  $< 1$ ; if different, loss  $> 1$ ; at boundary, loss=1
- This loss function is smooth & convex (easier to optimize)



# AdaBoost example: Viola-Jones


- Viola-Jones face detection algorithm
- Combine lots of very weak classifiers
  - Decision stumps = threshold on a single feature
- Define lots and lots of features
- Use AdaBoost to find good features
  - And weights for combining as well



# Haar wavelet features

- Four basic types.
  - They are easy to calculate.
  - The white areas are subtracted from the black ones.
  - A special representation of the sample called the **integral image** makes feature extraction faster.

Four types:




24x24 data : type, location, size => 162,336 features:




$x^{(i)}$




$\Phi_{18280}(\cdot)$



$\Phi_{126740}(\cdot)$



$\Phi_{9816}(\cdot)$



$\Phi_{36834}(\cdot)$

# Training a face detector

---

- Wavelets give ~100k features
- Each feature is one possible classifier
- To train: iterate from 1:T
  - Train a classifier on each feature using weights
  - Choose the best one, find errors and re-weight
- This can take a long time... (lots of classifiers)
  - One way to speed up is to not train very well...
  - Rely on adaboost to fix “even weaker” classifier
- Lots of other tricks in “real” Viola-Jones
  - Cascade of decisions instead of weighted combo
  - Apply at multiple image scales
  - Work to make computationally efficient

# Summary: Ensembles

---

- combine multiple trained models to make prediction
- Types
  - Committees: vote / average
  - Stacking: learn to combine
  - Mixtures of Experts: different “areas of expertise”
- Bagging
  - Randomly re-sample data to build diverse predictors
  - Averaging process reduces overfit of individual models
- Boosting
  - Train model to correct “remaining” mistakes
  - Combined model is more complex than individual models